

CS118 Discussion 1B, Week 7

Zhehui Zhang

HAINES A2 / Friday / 12:00pm-1:50pm

Outline

- Network data plane
 - Fragmentation, DHCP, NAT, IPv6, Openflow
- Network control plane
 - Routing
 - Link state routing
 - Distance vector routing
- Project 2

Quick question on fragmentation

- Consider following IP packet

4	5	TOS	2400			
12345			0	0	0	0
25		6	checksum			
10.1.1.1						
80.233.250.61						

- Assume MTU = 1450 Bytes. Show the header length, total length, identification, flags, fragment offset, TTL, and IP payload size.

Quick question

- Consider following IP packet

4	5	TOS	2400			
12345			0	0	0	0
25		6	checksum			
10.1.1.1						
80.233.250.61						

- Assume MTU = 1450 Bytes. Show the header length, total length, identification, flags, fragment offset, TTL, and IP payload size.

For the first packet: 5 (20 bytes), 1444 bytes, ID = 12345, MF = 1, Offset = 0, TTL = 25, 1424 bytes.
For the second packet: 5 (20 bytes), 976 bytes, ID = 12345, MF = 0, Offset = 178, TTL = 25, 956 bytes.

DHCP: Dynamic Host Configuration Protocol


- Dynamically allocates the following info to a host
 - IP address on subnet for the host
 - IP address for default router (“first-hop” router)
 - Subnet mask
 - IP address and name for DNS caching resolver
- Allows address reuse

DHCP: operations

- Host broadcasts “DHCP discovery” msg [optional]
- DHCP server responds with “DHCP offer” msg [optional]
- Host requests IP address: “DHCP request” msg
- DHCP server sends address: “DHCP ack” msg

Important example on Chapter 4 slides 45—46!

Quick question




- What information in DHCP will be updated if PC0 move to PC1?
- How about PC0 move to PC3?

NAT (network address translation)

- Depletion of IPv4 addresses — short-term solution
 - IP tunneling?
- Use private IP addresses
- Side-benefit: security
- How to achieve?
 - <public IP:port> — <private IP:port> mapping

Quick question

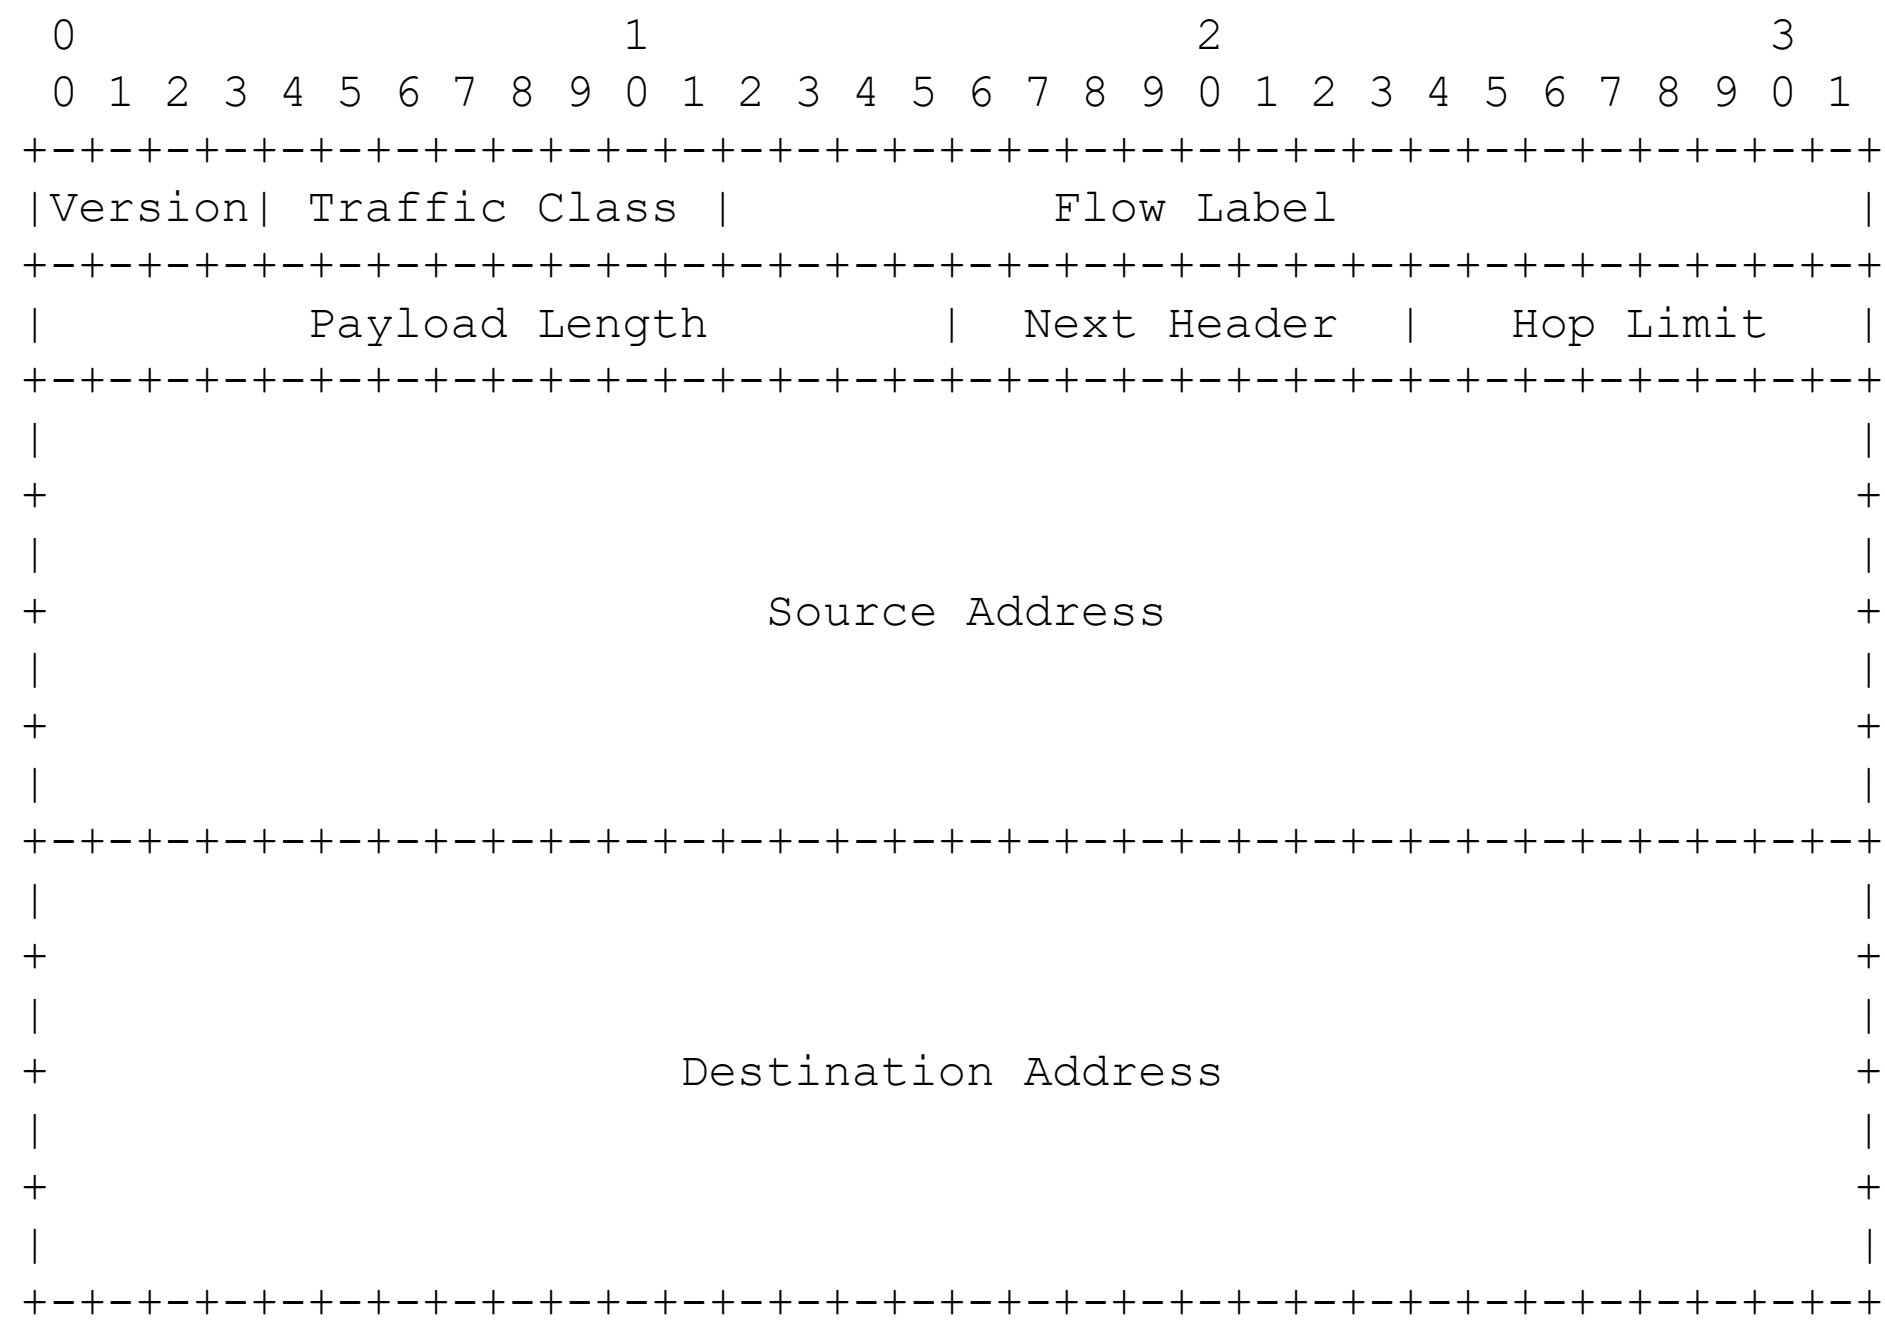
-



NAT: downside

- Increased complexity
- Single point of failure
- Cannot run services inside a NAT box
 - Why?

IPv6



IPv6 Header Format (RFC 2460)

IPv6/IPv4 differences

- Fixed-length 40 byte header
 - length field excludes header
 - Header Length field eliminated
- Address length: 128 bits
- Priority: usage yet to be finalized
- Flow Label: identify packets in same flow
- Next header: identify upper layer protocol for data
- Options: outside of the basic header, indicated by Next Header field
- Header Checksum: removed

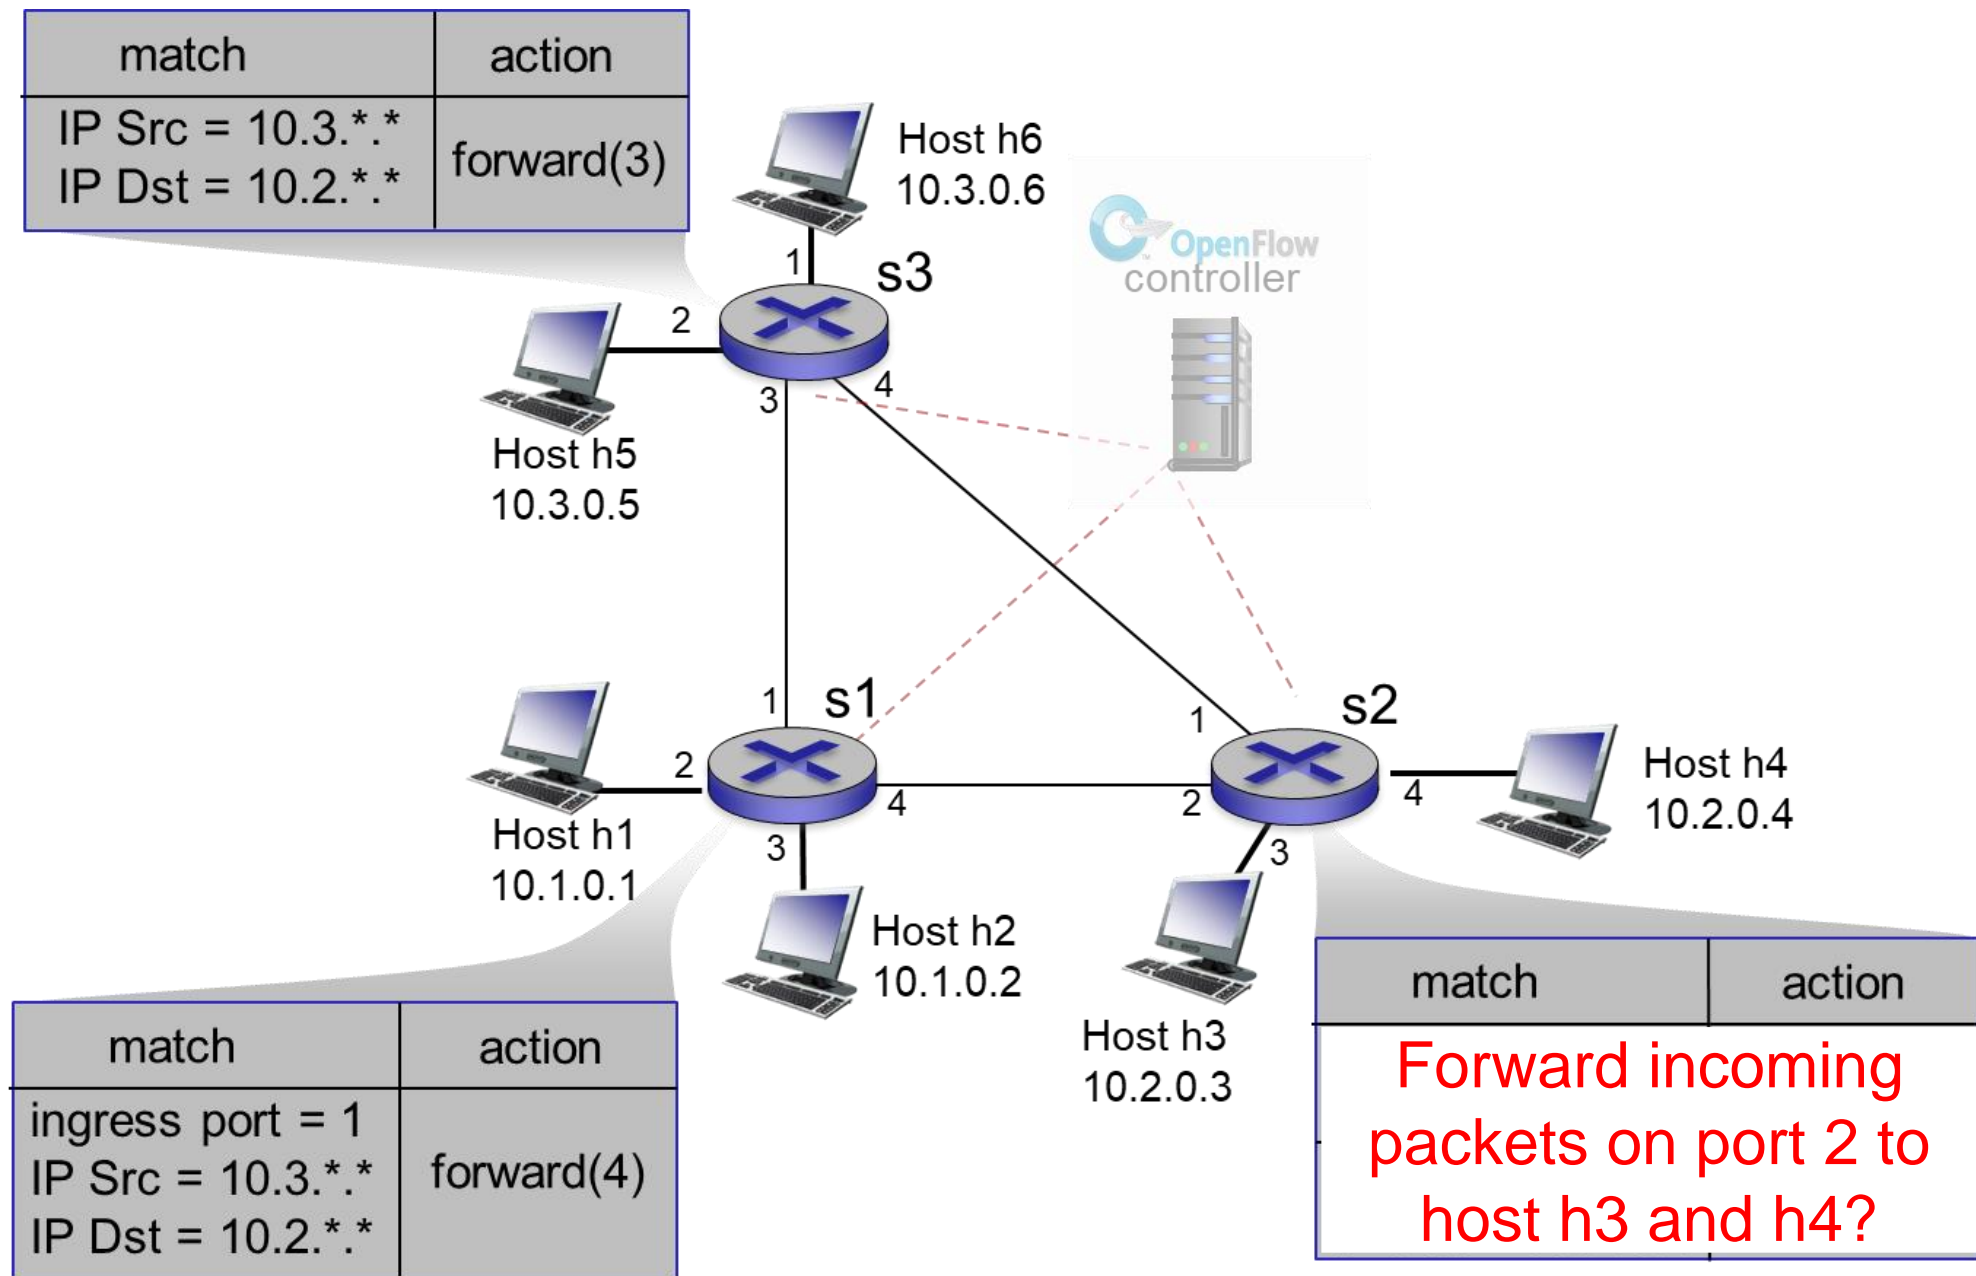
IPv6 address format (optional)

- Colon-Hex: 2607:F010:03f9:0000:0000:0000:0004:0001
- Can skip leading zeros of each word:
2607:F010:3f9:0:0:0:4:1
- Can skip one sequence of zero words (compressed representation), e.g., 2607:f010:3f9::4:1
- Can leave the last 32 bits in dot-decimal:
2607:f010:3f9::0.4.0.1
- Can specify a prefix by /length: 2607:f010:3f9::/64

Special IPv6 addresses (optional)

- `::/128` - Unspecified
- `::1/128` - Loopback
- `::ffff:0:0/96` - IP4-mapped address
- `2002::/16` - 6to4
- `ff00::/8` - Multicast
- `fe80::/10` - Link-Local Unicast

Open flow example



Routing: concepts

- Global or decentralized information?
 - global: all routers have complete topology, link cost info
 - algorithm?

Routing: concepts

- Global or decentralized information?
 - global: all routers have complete topology, link cost info
 - “link state” algorithms

Routing: concepts

- Global or decentralized information?
 - global: all routers have complete topology, link cost info
 - “link state” algorithms
 - decentralized: router knows physically-connected neighbors, link costs to neighbors; iterative process of computation, exchange of info with neighbors
 - algorithm?

Routing: concepts

- Global or decentralized information?
 - global: all routers have complete topology, link cost info
 - “link state” algorithms
 - decentralized: router knows physically-connected neighbors, link costs to neighbors; iterative process of computation, exchange of info with neighbors
 - “distance vector” algorithms

Link state routing

- Dijkstra's algorithm
 - net topology, link costs known to all nodes
 - computes least cost paths from one node ('source') to all other nodes
 - iterative: after k iterations, know least cost path to k destinations

Link state routing: algorithm

```
1  Initialization:
2    N' = {u}
3    for all nodes v
4      if v adjacent to u
5        then D(v) = c(u, v)
6      else D(v) = ∞
7
8  Loop
9    find w not in N' such that D(w) is a minimum
10   add w to N'
11   update D(v) for all v adjacent to w and not in N':
12     [Link cost update heuristic from Dijkstra algo.]
13  until all nodes in N'
```

$c(x, y)$: link cost from node x to y ; $c(x, y) = \infty$ if not direct neighbors

$D(v)$: current value of cost of path from source to destination v

$p(v)$: predecessor node along path from source to v

N' : set of nodes whose least cost path definitively known

Link state routing: algorithm

```
1  Initialization:
2     $N' = \{u\}$ 
3    for all nodes  $v$ 
4      if  $v$  adjacent to  $u$ 
5        then  $D(v) = c(u, v)$ 
6      else  $D(v) = \infty$ 
7
8  Loop
9    find  $w$  not in  $N'$  such that  $D(w)$  is a minimum
10   add  $w$  to  $N'$ 
11   update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $N'$ :
12      $D(v) = \min(D(v), D(w) + c(w, v))$ 
13  until all nodes in  $N'$ 
```

$c(x, y)$: link cost from node x to y ; $c(x, y) = \infty$ if not direct neighbors


$D(v)$: current value of cost of path from source to destination v

$p(v)$: predecessor node along path from source to v

N' : set of nodes whose least cost path definitively known

Link state routing: example

- Using link state routing to setup a forwarding table for node u



Let's work it out

N'	D(v), p(v)	D(w), p(w)	D(x), p(x)	D(y), p(y)	D(z), p(z)
u	2, u	5, u	1, u	∞	∞
ux	2, u	4, x		2, x	∞
uxy	2, u	3, y			4, y
uxyv		3, y			4, y
uxyvw					4, y
uxyvwz					

Let's work it out

N'	D(v), p(v)	D(w), p(w)	D(x), p(x)	D(y), p(y)	D(z), p(z)
u	2, u	5, u	1, u	∞	∞
ux	2, u	4, x		2, x	∞
uxy	2, u	3, y			4, y
uxyv		3, y			4, y
uxyvw					4, y
uxyvwz					

Link state routing: complexity

- size: n nodes
- each iteration: need to check all nodes, w , not in N
- $n(n+1)/2$ comparisons: $O(n^2)$
- more efficient implementations possible: $O(n \log n)$

Distance vector routing


- Bellman-Ford equation (dynamic programming)
- let
- $dx(y) := \text{cost of least-cost path from } x \text{ to } y$
- then
- $dx(y) = ?$

Distance vector routing

- Bellman-Ford equation (dynamic programming)
- let
- $dx(y) := \text{cost of least-cost path from } x \text{ to } y$
- then
- $dx(y) = \min_v \{c(x,v) + dv(y)\}, v: \text{neighbors of } x$


Distance vector routing: example

- What's the cost of least-cost path for $u \rightarrow z$?




Let's work it out

- clearly:
 - $dv(z) = ?$, $dx(z) = ?$, $dw(z) = ?$



Let's work it out

- clearly:
 - $dv(z) = 5$, $dx(z) = 3$, $dw(z) = 3$
- According to B-F equation:
 - $du(z) = \min \{ ? \}$



Let's work it out

- clearly:
 - $dv(z) = 5, dx(z) = 3, dw(z) = 3$
- According to B-F equation:
 - $du(z) = \min \{c(v, x) + dv(z), c(u, x) + dx(z), c(u, w) + dw(z)\}$

Let's work it out


- clearly:
 - $dv(z) = 5, dx(z) = 3, dw(z) = 3$
- According to B-F equation:
 - $du(z) = \min \{c(u, v) + dv(z), c(u, x) + dx(z), c(u, w) + dw(z)\}$
 - $= \min \{2 + 5, 1 + 3, 5 + 3\} = 4$

Distance vector routing: key idea

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when x receives new DV estimate from neighbor, it updates its own DV using B-F equation.


Distance vector routing: caveat

- Count-to-infinity problem.
- Can you work out an example?




Distance vector routing: caveat

- Count-to-infinity problem.
- Can you work out an example?



Distance vector routing: caveat


- Count-to-infinity problem.
- Can you work out an example?



Distance vector routing: caveat

- Count-to-infinity problem.
- Can you work out an example?
- Can you propose a solution?
- basic idea?

A should not propagate its distance to B!




Distance vector routing: split horizon

- Previous solution idea:
 - split horizon
 - if A reaches C through B, A should not tell B that B can reach C
 - Then B will not attempt to go through A to reach C
 - Are we good?

Distance vector routing: split horizon

- Previous solution idea:
 - split horizon
 - if A reaches C through B, A should not tell B that B can reach C
 - Then B will not attempt to go through A to reach C
 - Are we good?





Distance vector routing: poison reverse

- Split horizon + poison reverse
 - if A reaches D through C:
 - A tells C that A's distance to D is infinite
 - Then C will not attempt to go through A to reach D
 - In practice, infinite == 16 hops

Distance vector routing: poison reverse

If Z routes through Y to get to X:

- Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)



Link State v.s. Distance Vector

	Link state	Distance vector
message complexity	with n nodes, E links, $O(nE)$ msgs sent	exchange between neighbors only (convergence time varies)
convergence speed	$O(n^2)$ algorithm requires $O(nE)$ msgs	convergence time varies (may be routing loops)
robustness	node can advertise incorrect link cost; each node computes only its own table	DV node can advertise incorrect path cost; error propagate thru network
implementation	OSPF	RIP

Summary

- Link-state routing (Dijkstra) algorithm:
 - each node computes the shortest paths to all the other nodes based on the complete topology map
- Distance Vector (Bellman-Ford) routing algorithm:
 - each node computes the shortest paths to all the other nodes based on its neighbors distance to all destinations